

Neural Networks

CMPT 419/726

Mo Chen

SFU Computing Science

Jan. 29, 2020

Bishop PRML Ch. 5

Neural Networks

- Neural networks arise from attempts to model human/animal brains
 - Many models, many claims of biological plausibility
- We will focus on **multi-layer perceptrons**
 - Mathematical properties rather than plausibility



Applications of Neural Networks

- Many success stories for neural networks, old and new
 - Credit card fraud detection
 - Hand-written digit recognition
 - Face detection
 - Autonomous driving (CMU ALVINN)
 - Object recognition
 - Speech recognition

Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

Feed-forward Networks

- We have looked at generalized linear models of the form:

$$y(\mathbf{x}, \mathbf{w}) = f\left(\sum_{j=1}^M w_j \phi_j(\mathbf{x})\right)$$

for fixed non-linear basis functions $\phi(\cdot)$

- We now extend this model by allowing adaptive basis functions, and learning their parameters
- In feed-forward networks (a.k.a. **multi-layer perceptrons**) we let each basis function be another non-linear function of linear combination of the inputs:

$$\phi_j(\mathbf{x}) = f\left(\sum_{j=1}^M \dots\right)$$


Feed-forward Networks

- Starting with input $x = (x_1, \dots, x_D)$, construct linear combinations:

$$a_j = \sum_{i=1}^D (w_{ji}^{(1)} x_i + x_{j0}^{(1)})$$


These a_j are known as **activations**

- Pass through an **activation function** $h(\cdot)$ to get output $z_j = h(a_j)$
 - Model of an individual neuron



Activation Functions

- Can use a variety of activation functions
 - Sigmoidal (S-shaped)
 - Logistic sigmoid $1/(1 + \exp(-a))$ (useful for binary classification)
 - Hyperbolic tangent $\tanh(\cdot)$
 - Radial basis function $z_j = \sum_i (x_i - w_{ji})^2$
 - Softmax
 - Useful for multi-class classification
 - Identity
 - Useful for regression
 - Threshold
 - ...
- Needs to be differentiable for gradient-based learning (later)
- Can use different activation functions in each unit



Activation Functions

Common choices of activation functions

Softplus:

$$\log(1 + e^x)$$




Hyperbolic tangent:

$$\tanh x$$


Rectified linear unit (ReLU):

$$\max(0, x)$$

Key feature: easy to differentiate



Feed-forward Networks



- Connect together a number of these units into a feed-forward network (DAG)
- Above shows a network with one layer of **hidden units**
- Implements function:

$$y_k(x, w) = h \left(\sum_{j=1}^M w_{kj}^{(2)} h \left(\sum_{i=1}^D w_{ij}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right)$$

Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

Network Training

- Given a specified network structure, how do we set its parameters (weights)?
 - As usual, we define a criterion to measure how well our network performs, optimize against it
- For regression, training data are $(x_n, t_n), t_n \in \mathbb{R}$
 - Squared error naturally arises:


$$E(w) = \sum_{n=1}^N \{y(x_n, w) - t_n\}^2$$

- For binary classification, this is another discriminative model, ML:

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} \{1 - y_n\}^{1-t_n}$$


$$E(w) = - \sum_{n=1}^N \{t_n \ln(y_n) + (1 - t_n) \ln(1 - y_n)\}$$

Parameter Optimization



- For either of these problems, the error function $E(\mathbf{w})$ is nasty
 - Nasty = non-convex
 - Non-convex = has **local minima**


A Non-Convex function



Aside: Optimization Program

$$\begin{array}{ll}\text{minimize} & f(x) \\ \text{subject to} & g_i(x) \leq 0, i = 1, \dots, n \\ & h_j(x) = 0, j = 1, \dots, m\end{array}$$

- Very difficult to solve in general
 - Trade-offs to consider: computation time, solution optimality
- Easy cases:
 - Find global optimum for **linear program**: f, g_i, h_j are linear
 - Find global optimum for **convex program**: f, g_i are convex, h_j is linear
 - Find local optimum for **nonlinear program**: f, g_i, h_j are differentiable
- Neural Networks: Nonlinear and unconstrained



Convex Functions



- **Convex function**

$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y)$ for all $x, y \in \mathbb{R}^n$, for all $\theta \in [0, 1]$

- Sublevel sets of convex functions, $\{x: f(x) \leq C\}$, are convex

- **Convex shape \mathcal{C} :**


$x_1, x_2 \in \mathcal{C}, \theta \in [0, 1] \Rightarrow \theta x_1 + (1 - \theta)x_2 \in \mathcal{C}$



Convex Functions

- **Convex function**

$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y)$ for all $x, y \in \mathbb{R}^n$, for all $\theta \in [0,1]$




- Sublevel sets of convex functions, $\{x: f(x) \leq C\}$, are convex

- **Convex shape \mathcal{C} :**

$x_1, x_2 \in \mathcal{C}, \theta \in [0,1] \Rightarrow \theta x_1 + (1 - \theta)x_2 \in \mathcal{C}$


- Superlevel sets of convex functions are *not* convex!





Common Convex Functions on \mathbb{R}

- $f(x) = e^{ax}$ is convex for all $x, a \in \mathbb{R}$
- $f(x) = x^a$ is convex on $x > 0$ if $a \geq 1$ or $a \leq 0$; concave if $0 < a < 1$
- $f(x) = \log x$ is concave
- $f(x) = x \log x$ is convex for $x > 0$ (or $x \geq 0$ if defined to be 0 when $x = 0$)

$$f(x) = e^{ax}$$




$$f(x) = x^a$$





Common Convex Functions on \mathbb{R}^n

- $f(x) = Ax + b$ is convex for any A, b
- Every norm on \mathbb{R}^n is convex
- $f(x) = \max(x_1, x_2, \dots, x_n)$ is convex
- $f(x) = \frac{x_1^2}{x_2}$ (for $x_2 > 0$)
- Log-sum-exp softmax: $f(x) = \frac{1}{k} \log(e^{kx_1} + e^{kx_2} + \dots + e^{kx_n})$
- Geometric mean: $f(x) = (\prod_{i=1}^n x_i)^{\frac{1}{n}}, x_i > 0$



$$f(x) = \frac{1}{5} \log(e^{5x_1} + e^{5x_2})$$



Descent Methods



- The typical strategy for optimization problems of this sort is a descent method:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$

- As we've seen before, these come in many flavours
 - Gradient descent $\nabla E(\mathbf{w}^{(\tau)})$
 - Stochastic gradient descent $\nabla E_n(\mathbf{w}^{(\tau)})$
 - Newton-Raphson (second order)
- All of these can be used here, stochastic gradient descent is particularly effective
 - Redundancy in training data, escaping local minima


Numerical Solution: Gradient Methods

- Start from x^0 and construct a sequence x^k such that $x^k \rightarrow x^*$
 - Calculate x^{k+1} from x^k by “going down the gradient”
 - Unconstrained case: $x^{k+1} = x^k - \alpha^k \nabla f(x)$, $\alpha^k > 0$




Numerical Solution: Gradient Methods

- Start from x^0 and construct a sequence x^k such that $x^k \rightarrow x^*$
 - Calculate x^{k+1} from x^k by “going down the gradient”
 - Unconstrained case: $x^{k+1} = x^k - \alpha^k \nabla f(x)$, $\alpha^k > 0$
- More generally, $x^{k+1} = x^k + \alpha^k d^k$ for some d such that
$$\nabla f(x^k) \cdot d^k < 0$$
- Tuning parameters: descent direction d^k , and step size α^k




Descent Direction

- Steepest descent: $d^k = -\nabla f(x^k)$
 - $x^{k+1} = x^k - \alpha^k \nabla f(x)$
 - Simple but sometimes leads to slow convergence
- Newton's method: $d^k = \left(\nabla^2 f(x^k)\right)^{-1} \nabla f(x^k)$
 - Minimize the quadratic approximation:
$$f^k(x) = f(x^k) + \nabla f(x^k)^\top (x - x^k) + \frac{1}{2}(x - x^k)^\top \nabla^2 f(x^k)(x - x^k)$$
 - Set gradient to zero to obtain next iterate
$$\begin{aligned}\nabla f^k(x) &= \nabla f(x^k) + \nabla^2 f(x^k)(x - x^k) = 0 \\ \Rightarrow x^{k+1} &= x^k - \left(\nabla^2 f(x^k)\right)^{-1} \nabla f(x^k)\end{aligned}$$
 - Fast convergence, but matrix inverse required
 - Alternatively, use an algorithm to minimize a quadratic function



Step Size

- Recall $x^{k+1} = x^k + \alpha^k d^k$, with $\nabla f(x^k)^\top d^k < 0$
- Line search: choose $\alpha^k = \min_{\alpha \geq 0} f(x^k + \alpha^k d^k)$
 - Requires minimization
- Constant step size: $\alpha^k = \alpha$
 - May not converge
- Diminishing step size: $\alpha^k \rightarrow 0$
 - Still need to explore all regions $\sum \alpha^k = \infty$
 - For example: $\alpha^k = \frac{\alpha^0}{k}$



Numerical Solution: Second Order Methods

$$\text{minimize } f(x) \quad \longrightarrow \quad \text{minimize }_{d_x} (\mathbf{r}^k)^\top d_x + \frac{1}{2} d_x^\top \mathbf{B}_k d_x$$



where $d_x := x - x^k$,

- Quadratize $f(x)$:

$$\mathbf{r}^k = \nabla f(x_k)$$


$$\mathbf{B}_k = \mathbf{H}f(x_k)$$

- Convexify if needed, eg. by removing negative eigenvalues



Example

minimize $0.5x^4 + 0.8x^3 - 3x^2 - 2x + 5$
subject to $-3 \leq x \leq 2$



Computing Gradients

- The function $y(\mathbf{x}_n, \mathbf{w})$ implemented by a network is complicated
 - It isn't obvious how to compute error function derivatives with respect to weights
- Numerical method for calculating error derivatives, use finite differences:

$$\frac{\partial E_n}{\partial w_{ji}} \approx \frac{E_n(w_{ji} + \epsilon) - E_n(w_{ji} - \epsilon)}{2\epsilon}$$

- How much computation would this take with W weights in the network?
 - $O(|W|)$ per partial derivative (evaluation of E_n)
 - $O(|W|^2)$ total per gradient descent step (there are $|W|$ partial derivatives)

Outline


Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

Feed-forward Networks



- Connect together a number of these units into a feed-forward network (DAG)
- Above shows a network with one layer of **hidden units**
- Implements function:

$$y_{(n),k}(x_n, w) = h \left(\sum_{j=1}^M w_{kj}^{(2)} \underbrace{h \left(\sum_{i=1}^D w_{ji}^{(1)} x_{(n),i} + w_{j0}^{(1)} \right)}_{z_{(n),j}} + w_{k0}^{(2)} \right)$$

Error Backpropagation

- Backprop is an efficient method for computing error derivatives $\frac{\partial E_n}{\partial w_{ji}^{(m)}}$
 - $O(W)$ to compute derivatives wrt all weights
- First, feed training example x_n forward through the network, storing all activations a_j
- Calculating derivatives for weights connected to output nodes is easy

- e.g. For linear output nodes $y_k = \sum_i w_{ki}^{(L)} z_{(n),i}^{(L-1)}$:

$$\frac{\partial E_n}{\partial w_{ki}^{(L)}} = \frac{\partial}{\partial w_{ki}^{(L)}} \frac{1}{2} (y_{(n),k} - t_{(n),k})^2 = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)}$$

- For hidden layers, propagate error backwards from the output nodes

Error Backpropagation

$y_{(n),k}, E_n$:

- n : data point
- k : component


$w_{ji}^{(m)}$:

- m : layer
- j : index matching output
- i : index matching input

$$E(w) = \frac{1}{2} \sum_{n=1}^N \sum_k (y_{(n),k} - t_{(n),k})^2, \quad y_{(n),k} = \sum_i w_{ki}^{(L)} z_{(n),i}^{(L-1)}$$

$$E_n(w) = \frac{1}{2} \sum_k (y_{(n),k} - t_{(n),k})^2$$

$$\frac{\partial E_n}{\partial w_{ki}^{(L)}} = \frac{\partial}{\partial w_{ki}^{(L)}} \frac{1}{2} \sum_{k'} (y_{(n),k'} - t_{(n),k'})^2 = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)} \quad (*)$$



Chain Rule for Partial Derivatives

- A “reminder”
- For $f(x, y)$, with f differentiable wrt x and y , and x and y differentiable wrt u :

$$\frac{\partial f}{\partial u} = \frac{\partial f}{\partial x} \frac{\partial x}{\partial u} + \frac{\partial f}{\partial y} \frac{\partial y}{\partial u}$$


Error Backpropagation

- We can write

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \frac{\partial}{\partial w_{ji}^{(m)}} E_n \left(a_{(n),1}^{(m)}, a_{(n),2}^{(m)}, \dots, a_{(n),D}^{(m)} \right)$$

- Using the chain rule:

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} \frac{\partial a_{(n),1}^{(m)}}{\partial w_{ji}^{(m)}} + \sum_{k \neq j} \frac{\partial E_n}{\partial a_{(n),k}^{(m)}} \frac{\partial a_{(n),k}^{(m)}}{\partial w_{ji}^{(m)}}$$



where $\sum_k(\dots)$ runs over all other nodes k in the same layer (m)

- Since $a_{(n),k}^{(m)}$ does not depend on $w_{ji}^{(m)}$, all terms in the summation go to 0:

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} \frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}}$$

Error Backpropagation cont.


- Introduce error $\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}}$

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \delta_{(n),j}^{(m)} \frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}}$$

- Other factor is

$$\frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}} = \frac{\partial}{\partial w_{ji}^{(m)}} \sum_k w_{jk}^{(m)} z_k^{(m-1)} = z_i^{(m-1)}$$

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \delta_{(n),j}^{(m)} z_i^{(m-1)}$$



Error Backpropagation cont.

- Error $\delta_{(n),j}^{(m)}$ can also be computed using chain rule:


$$\delta_j^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} = \sum_k \underbrace{\frac{\partial E_n}{\partial a_{(n),k}^{(m+1)}}}_{\delta_k^{(m+1)}} \frac{\partial a_{(n),k}^{(m+1)}}{\partial a_{(n),j}^{(m)}}$$

where $\sum_k(\dots)$ runs over all nodes k in the layer **after**.

$$a_{(n),k}^{(m+1)} = \sum_i w_{ki}^{(m+1)} z_{(n),i}^{(m)} = \sum_i w_{ki}^{(m+1)} h^{(m)}(a_{(n),i}^{(m)})$$

$$\frac{\partial a_{(n),k}^{(m+1)}}{\partial a_{(n),j}^{(m)}} = w_{kj}^{(m+1)} (h^{(m)})' (a_{(n),j}^{(m)})$$

$$\delta_{(n),j}^{(m)} = \sum_k \delta_{(n),k}^{(m+1)} w_{kj}^{(m+1)} (h^{(m)})' (a_{(n),j}^{(m)}) = (h^{(m)})' (a_{(n),j}^{(m)}) \sum_k \delta_{(n),k}^{(m+1)} w_{kj}^{(m+1)}$$



Error Backpropagation cont.

- **Error** $\delta_{(n),j}^{(m)}$ can also be computed using chain rule:

$$\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} = \sum_k \underbrace{\frac{\partial E_n}{\partial a_{(n),k}^{(m+1)}}}_{\delta_k} \frac{\partial a_{(n),k}^{(m+1)}}{\partial a_{(n),j}^{(m)}}$$

where $\sum_k(\cdots)$ runs over all nodes k in the layer **after**.

- Eventually:

$$\delta_{(n),j}^{(m)} = (h^{(m)})' (a_{(n),j}^{(m)}) \sum_k \delta_{(n),k}^{(m+1)} w_{kj}^{(m+1)}$$

- A weighted sum of the later error “caused” by this weight

Error Backpropagation cont.

- Eventually:

$$\delta_{(n),j}^{(m)} = (h^{(m)})' (a_{(n),j}^{(m)}) \sum_k \delta_{(n),k}^{(m+1)} w_{jk}^{(m+1)}$$

where $\sum_k(\dots)$ runs over all nodes k in the layer **after**.

- Above recursion relation needs last set of errors: $\delta_j^{(L)}$

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \delta_{(n),j}^{(m)} z_i^{(m-1)} \quad (\text{by definition})$$

$$\frac{\partial E_n}{\partial w_{ji}^{(L)}} = \delta_{(n),j}^{(L)} z_{(n),i}^{(L-1)} = (y_{(n),j} - t_{(n),j}) z_{(n),i}^{(L-1)} \quad (\text{from before } (*))$$

$$\delta_{(n),j}^{(L)} = y_{(n),j} - t_{(n),j} \quad (\text{by comparison})$$

Summary

 $O(|W|)$

Output Definition / forward propagation

$$y_{(n),k}(x_n, w) = h^{(m+1)} \left(\sum_{j=1}^M w_{jk}^{(m+1)} h^{(m)} \left(\sum_{i=1}^D w_{ij}^{(m)} z_{(n),i}^{(m-1)} + w_{0j}^{(m)} \right) + w_{k0}^{(m+1)} \right)$$

- Save \mathbf{z}, \mathbf{a}

Gradient computation / backpropagation


- Last layer: $\frac{\partial E_n}{\partial w_{ik}^{(L)}} = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)}$
- Previous layers: Define $\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}}$

Starting from last layer,

$$\delta_{(n),j}^{(L)} = y_{(n),j} - t_{(n),j}$$

Recursion: $\frac{\partial E_n}{\partial w_{ij}^{(m)}} = \delta_{(n),j}^{(m)} z_{(n),i}^{(m-1)}$,

where $\delta_{(n),j}^{(m)} = (h^{(m)})' \left(a_{(n),j}^{(m)} \right) \sum_k \delta_k^{(m+1)} w_{jk}^{(m+1)}$



Summary

Output Definition / forward propagation

$$y_{(n),k}(x_n, w) = h^{(m+1)} \left(\sum_{j=1}^M w_{jk}^{(m+1)} h^{(m)} \left(\sum_{i=1}^D w_{ij}^{(m)} z_{(n),i}^{(m-1)} + w_{0j}^{(m)} \right) + w_{k0}^{(m+1)} \right)$$

$O(|W|)$

- Save \mathbf{z} , \mathbf{a}

Gradient computation / backpropagation

- Last layer: $\frac{\partial E_n}{\partial w_{ik}^{(L)}} = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)}$
- Previous layers: Define $\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}}$

Starting from last layer,

$$\delta_{(n),j}^{(L)} = y_{(n),j} - t_{(n),j}$$

Recursion: $\frac{\partial E_n}{\partial w_{ij}^{(m)}} = \delta_{(n),j}^{(m)} z_{(n),i}^{(m-1)}$,

where $\delta_{(n),j}^{(m)} = (h^{(m)})' \left(a_{(n),j}^{(m)} \right) \sum_k \delta_k^{(m+1)} w_{jk}^{(m+1)}$

Goes through one layer of weights



Goes through one layer of weights

Tensorflow Playground

- <https://playground.tensorflow.org>

Outline

Feed-forward Networks

Network Training


Error Backpropagation

Deep Learning

Deep Learning


- Collection of important techniques to improve performance:
 - Multi-layer networks
 - Convolutional networks, parameter tying
 - Hinge activation functions (ReLU) for steeper gradients
 - Momentum
 - Drop-out regularization
 - Sparsity
 - Auto-encoders for unsupervised feature learning
 - ...
- **Scalability** is key, can use lots of data since stochastic gradient descent is memory-efficient, can be parallelized

Hand-written Digit Recognition




- MNIST - standard dataset for hand-written digit recognition
 - 60000 training, 10000 test images

LeNet-5, circa 1998



- LeNet developed by Yann LeCun et al.
 - Convolutional neural network
 - Local receptive fields (5x5 connectivity)
 - Subsampling (2x2)
 - Shared weights (reuse same 5x5 “filter”)
 - Breaking symmetry

ImageNet



- ImageNet - standard dataset for object recognition in images (Russakovsky et al.)
 - 1000 image categories, ≈ 1.2 million training images (ILSVRC 2013)

GoogLeNet, circa 2014




- GoogLeNet developed by Szegedy et al., CVPR 2015
- Modern deep network
- ImageNet top-5 error rate of 6.67% (later versions even better)
- Comparable to human performance (especially for fine-grained categories)

ResNet, circa 2015

- ResNet developed by He et al., ICCV 2015
- 152 layers
- ImageNet top-5 error rate of 3.57%
- Better than human performance (especially for fine-grained categories)



Key Component 1: Convolutional Filters




- Share parameters across network
- Reduce total number of parameters
- Provide **translation invariance**, useful for visual recognition

Common Operations

- Fully connected (dot product)
- Convolution
 - Translationally invariant
 - Controls overfitting
- Pooling (fixed function)
 - Down-sampling
 - Controls overfitting
- Nonlinearity layer (fixed function)
 - Activation functions, e.g. ReLU




Example: Small VGG Net From Stanford CS231n



Neural Network Architectures

- Convolutional neural network (CNN)

- Has translational invariance properties from convolution
- Common used for computer vision




- Recurrent neural network RNN

- Has feedback loops to capture temporal or sequential information
- Useful for handwriting recognition, speech recognition, reinforcement learning
- Long short-term memory (LSTM): special type of RNN with advantages in numerical properties

- Others


- General feedforward networks, variational autoencoders (VAEs), conditional VAEs



Training Neural Networks


- Data preprocessing
 - Removing bad data
 - Transform input data (e.g. rotating, stretching, adding noise)
- Training process (optimization algorithm)
 - Choice of loss function (eg. L1 and L2 regularization)
 - Dropout: randomly set neurons to zero in each training iteration
 - **Learning rate** (step size) and other hyperparameter tuning
- Software packages: efficient gradient computation
 - Caffe, Torch, Theano, TensorFlow

Key Component 2: Rectified Linear Units (ReLUs)




- **Vanishing gradient** problem
 - If derivatives very small, no/little progress via stochastic gradient descent
 - Occurs with sigmoid function when activation is large in absolute value
- ReLU: $h(a_j) = \max(0, a_j)$
- Non-saturating, linear gradients (as long as non-negative activation on some training data)
- Sparsity inducing

Key Component 3: Many, Many Layers



- **ResNet**: ≈ 152 layers (“shortcut connections”)
- GoogLeNet: ≈ 27 layers (“Inception” modules)
- VGG Net: 16-19 layers (Simonyan and Zisserman, 2014)
- Supervision: 8 layers (Krizhevsky et al., 2012)


Key Component 3: Many, Many Layers



- ResNet: ≈ 152 layers (“shortcut connections”)
- **GoogLeNet**: ≈ 27 layers (“Inception” modules)
- VGG Net: 16-19 layers (Simonyan and Zisserman, 2014)
- Supervision: 8 layers (Krizhevsky et al., 2012)

Key Component 3: Many, Many Layers

- ResNet: ≈ 152 layers (“shortcut connections”)
- GoogLeNet: ≈ 27 layers (“Inception” modules)
- **VGG Net**: 16-19 layers (Simonyan and Zisserman, 2014)
- Supervision: 8 layers (Krizhevsky et al., 2012)



Key Component 4: Momentum


- Trick to escape plateaus / local minima
- Take exponential average of previous gradients

$$\frac{\overline{\partial E_n}^\tau}{\partial w_{ji}} = \frac{\overline{\partial E_n}^\tau}{\partial w_{ji}} + \alpha \frac{\overline{\partial E_n}^{\tau-1}}{\partial w_{ji}}$$


- Maintains progress in previous direction

Key Component 5: Asynchronous Stochastic Gradient Descent

- Big models won't fit in memory
- Want to use compute clusters (e.g. 1000s of machines) to run stochastic gradient descent
- How to parallelize computation?
- **Ignore synchronization across machines**
 - Just let each machine compute its own gradients and pass to a server storing current parameters
 - Ignore the fact that these updates are inconsistent
 - Seems to just work (e.g. Dean et al. NIPS 2012)



Key Component 6: Learning Rate Schedule



- How to set learning rate η ?:

$$\mathbf{w}^\tau = \mathbf{w}^{\tau-1} + \eta \nabla \mathbf{w}$$

- **Option 1:** Run until validation error plateaus. Drop learning rate by x%
- **Option 2:** Adagrad, adaptive gradient. Per-element learning rate set based on local geometry (Duchi et al. 2010)

Key Component 7: Data Augmentation



- Augment data with additional synthetic variants (10x amount of data)
- Or just use synthetic data, e.g. Sintel animated movie (Butler et al. 2012)

Key Component 8: Data and Compute



- Get lots of data (e.g. ImageNet)
- Get lots of compute (e.g. CPU cluster, GPUs)
- Cross-validate like crazy, train models for 2-3 weeks on a GPU
- **Researcher gradient descent (RGD)** or **Graduate student descent (GSD)**: get 100s of researchers to each do this, trying different network structures

Challenges

Interpretability:



"panda"

57.7% confidence

+ ϵ



=



"gibbon"

99.3% confidence

Challenges

Data efficiency:

- ImageNet: 14 million images, 20000 categories
- AlphaStar: 200 years of gameplay



Challenges

- Problem formulation (what are you trying to predict?)
- Choice of model and optimization algorithm
- Data collection, post-processing
- Feature selection
- ...

More information

- <https://sites.google.com/site/deeplearningsummerschool>
- <http://tutorial.caffe.berkeleyvision.org/>
- ufldl.stanford.edu/eccv10-tutorial
- <http://www.image-net.org/challenges/LSVRC/2012/supervision.pdf>
- **Prof. Oliver Schulte's CMPT880: Deep Learning**
- **Project ideas**
 - Long short-term memory (LSTM) models for temporal data
 - Learning embeddings (word2vec, FaceNet)
 - Structured output (multiple outputs from a network)
 - Zero-shot learning (learning to recognize new concepts without training data)
 - Transfer learning (use data from one domain/task, adapt to another)

Conclusion

- Readings: Ch. 5.1, 5.2, 5.3
- Feed-forward networks can be used for regression or classification
 - Similar to linear models, except with **adaptive** non-linear basis functions
 - These allow us to do more than e.g. linear decision boundaries
- Different error functions
- Learning is more difficult, error function not convex
 - Use stochastic gradient descent, obtain (good?) local minimum
- Backpropagation for efficient gradient computation